

Review of Class-Incremental Methods

Giulia Ghisolfi, g.ghisolfi@studenti.unipi.it, 664222

Master Degree in Computer in Science (AI Curriculum), University of Pisa

Continual Learning course (791AA)

Academic Year: 2023-2024

Introduction

Problem Setting

Methodology: Class-Incremental Classifier

Methodology: Generative Replay

Analysis and Comparison between Methods

Strengths and Weaknesses

Introduction

Incremental Learning is a paradigm where models learn new tasks continuously without retraining from scratch, reducing computational costs and reusing resources.

In **Class-Incremental Learning**, new classes are introduced incrementally at each step, older classes are not revisited, and algorithms must classify all classes encountered across time steps.

1. **Memory-based and rebalancing classifiers:** These methods rely on memory storage or rebalancing techniques to mitigate catastrophic forgetting and address the class imbalance.
 - **iCaRL** (Incremental Classifier and Representation Learning)
 - **IL2M** (Class-Incremental Learning with Dual Memory)
 - Learning a Unified Classifier Incrementally via Rebalancing
2. **Generative approaches:** These methods represent an alternative to memory-based strategies. Instead of storing real data, they save a generative model capable of synthesizing examples of older classes during training.
 - Memory Replay GANs

Problem Setting

Catastrophic Forgetting: New class learning overwrites prior knowledge, requiring regularization or loss approximation to preserve past classes.

Stability-Plasticity Trade-Off: Balancing retention of old knowledge (stability) and adaptation to new classes (plasticity).

Class Imbalance: Recent classes dominate predictions, leading to *classifier bias*.

Technique to transfers knowledge from a larger (*teacher*) to a smaller (*student*) model.

Mitigates forgetting by applying distillation to old classes and Cross-Entropy loss to all classes, addressing mismatched output distributions as the classifier grows.

Methodology: Class-Incremental Classifier

Combines knowledge representation and knowledge distillation to address catastrophic forgetting, class imbalance, and effectively manage a large number of classes over time.

Prototype-based Strategy (Nearest-Mean-of-Exemplars): Classifies inputs by finding the nearest class prototype in feature space.

- Class prototype (μ_y) is the mean feature vector of exemplars:

$$\mu_y = \frac{1}{|P_y|} \sum_{p \in P_y} \phi(p).$$

- Classify x by finding the nearest prototype:

$$y^* = \arg \min_{y=1,\dots,t} \|\phi(x) - \mu_y\|.$$

Incremental Classifier and Representation Learning

Representation Learning with Knowledge Distillation: Combines new and exemplar data to preserve old class knowledge while learning new ones.

- **Classification Loss:** Encourages accurate predictions for new classes.
- **Distillation Loss:** Preserves old class knowledge.

$$\mathcal{L}(\Theta) = \mathcal{L}_{\text{classification}} + \mathcal{L}_{\text{distillation}}.$$

Exemplar Selection (Herdin): Maintains a fixed memory size by selecting the most representative samples for each class.

- Allocates $m = \frac{K}{t}$ exemplars per class.
- For new classes, selects exemplars to minimize deviation from class mean: $p_k = \arg \min_{x \in X} \left\| \mu - \frac{1}{k} \left(\phi(x) + \sum_{j=1}^{k-1} \phi(p_j) \right) \right\|.$

Class-Incremental Learning with Dual Memory

Builds on the idea that classes are best modeled when first learned with all data available.

Introduces a dual memory structure:

- **First Memory:** Stores exemplar images of past classes as a representative subset.
- **Second Memory:** Stores initial class statistics, providing optimal representations with minimal storage cost.

Employs a distillation-based loss to align predictions and mitigate catastrophic forgetting.

Class-Incremental Learning with Dual Memory

Bounded Memory and Data Imbalance:

- Retains $m = \frac{K}{t}$ exemplars per class.
- Imbalance increases as classes grow.

Prediction Rectification: To correct biases toward new classes, predictions for past classes C_i are rectified using:

$$p^r(x) = \begin{cases} p(C_i) \times \frac{\mu(M_N)}{\mu(M_P)} \times \frac{\mu_N(C_i)}{\mu_P(C_i)}, & \text{if pred} = \text{new}, \\ p(C_i), & \text{otherwise.} \end{cases}$$

Efficiency: The rectification formula adds minimal computational overhead, maintaining the method's efficiency while reducing bias toward new classes.

Incrementally learn a unified classifier that systematically investigates and addresses the effects of class imbalance, improving training balance and reducing bias toward new classes.

Cosine Normalization: Ensures balanced predictions with normalized weight vectors:

$$p(x) = \frac{\exp(\eta \cdot \langle \bar{\theta}_i, \bar{f}(x) \rangle)}{\sum_j \exp(\eta \cdot \langle \bar{\theta}_j, \bar{f}(x) \rangle)}$$

Learning a Unified Classifier Incrementally via Rebalancing

Less-forget Constraint: Preserves the geometric configuration of old classes using distillation loss:

$$L_{LF} = 1 - \langle \bar{f}^*(x), \bar{f}(x) \rangle,$$

weighted by $\lambda = \lambda_{\text{base}} \cdot \frac{|C_n|}{|C_o|}$.

Inter-class Separation: Enhances separation between old and new classes using margin-based loss:

$$L_{\text{margin}} = \sum_{k=1}^K \max(0, m - \langle \bar{\theta}(x), \bar{f}(x) \rangle + \langle \bar{\theta}_k, \bar{f}(x) \rangle).$$

Final Loss: $L = \frac{1}{|N|} \sum_{x \in N} (L_{\text{ce}}(x) + \lambda L_{LF}(x)) + \frac{1}{|N_o|} \sum_{x \in N_o} L_{\text{margin}}(x).$

Methodology: Generative Replay

Address sequential learning by incorporating a conditional GAN framework with three components:

- **Generator (G):** Synthesizes data for both old and new classes.
- **Discriminator (D):** Distinguishes real from generated images.
- **Classifier (C):** Predicts labels, forcing G to generate meaningful images aligned with target labels.

Data conditioned on class label c and latent vector z : $\tilde{x} = G(z, c)$.

Model addresses both adversarial and classification tasks during training.

Memory Replay GANs

Sequential Learning in GANs: GAN training extended for incremental learning with parameters initialized from the previous task:

$$\min_{\theta_G^t} \mathcal{L}_{\text{GAN}}^G(\theta_t, S_t), \quad \min_{\theta_D^t} \mathcal{L}_{\text{GAN}}^D(\theta_t, S_t),$$

where $\theta_t = (\theta_G^t, \theta_D^t)$.

Elastic Weight Consolidation (EWC): Penalizes significant changes in important parameters to prevent forgetting:

$$\min_{\theta_G^t} \mathcal{L}_{\text{GAN}}(\theta_G^t, S_t) + \frac{\lambda_{\text{EWC}}}{2} \sum_i F_{t-1,i} (\theta_{G,t,i} - \theta_{G,t-1,i})^2.$$

Memory Replay Mechanism: *Joint Retraining with Replay:* Combines synthetic data from old tasks with real data from the current task for training. *Replay Alignment:* Synchronizes G_t and G_{t-1} to generate identical outputs for a given z, c .

Analysis and Comparison between Methods

Class-Incremental Learning Classifiers Comparison

Common Features: All methods use memory buffers, distillation techniques, and address class imbalance.

Differences:

- **iCaRL:** Nearest-mean classifier with fixed memory. Does not explicitly handle class imbalance, leading to lower accuracy for old classes.
- **IL2M:** Dual memory (exemplars + class statistics). Rectifies old class predictions, improving accuracy with low memory usage.
- **LUCIR:** Combines advanced techniques (cosine normalization, less-forget constraint, inter-class separation), to explicitly balance training and reduce bias toward new classes.

iCaRL vs IL2M: IL2M consistently outperforms iCaRL across all memory sizes and datasets (ILSVRC, VGGFace2, Landmarks), demonstrating superior accuracy and robustness, particularly with limited memory.

iCaRL vs LUCIR: The Unified Classifier improves accuracy by 6% on CIFAR-100 and reduces error by 13% on ImageNet compared to iCaRL, effectively balancing training and reducing bias toward new classes.

MeRGANs reduce catastrophic forgetting by generating high-quality samples of old classes, providing a scalable, memory-efficient alternative for class-incremental learning. However, further improvements are needed for complex datasets like CIFAR-100.

Authors also demonstrate that the generated images provide a visual representation of task interference and potential forgetting.

Strengths and Weaknesses

Strengths and Weaknesses

| Method | Strengths | Weaknesses |
|----------------|--|--|
| iCaRL | Simple and effective nearest-mean classifier. | Lacks mechanisms for class imbalance. |
| | Knowledge Distillation preserves knowledge of old classes. | Limited flexibility due to fixed-size memory. |
| IL2M | Two-memory structure corrects predictions. | Limited scalability due to compact statistics. |
| | Low computational overhead. | No explicit optimization of feature space. |
| | Strong theoretical justification for statistics. | |
| LUCIR | Combines techniques to address imbalance and forgetting. | Higher computational complexity. |
| | Explicitly balances training. | Requires careful hyperparameter tuning. |
| MeRGANs | Generative replay reduces memory requirements. | Instability during GAN training. |
| | Maintains old/new class classification. | Computationally expensive. |
| | Visual representation of task interference and forgetting. | Does not address feature space separation. |



E. Belouadah and A. Popescu.

II2m: Class incremental learning with dual memory.

In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 583–592, 2019.



S. Hou, X. Pan, C. C. Loy, Z. Wang, and D. Lin.

Learning a unified classifier incrementally via rebalancing.

In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 831–839, 2019.



S.-A. Rebuffi, A. Kolesnikov, G. Sperl, and C. H. Lampert.

icarl: Incremental classifier and representation learning.

In *Proceedings of the IEEE conference on Computer Vision and Pattern Recognition*, pages 2001–2010, 2017.



C. Wu, L. Herranz, X. Liu, J. Van De Weijer, B. Raducanu, et al.

Memory replay gans: Learning to generate new categories without forgetting.

Advances in neural information processing systems, 31, 2018.